✓ **Congratulations! You passed!**

Grade received 100%   **Latest Submission Grade** 100%   **To pass** 80% or higher

[**Go to next item**]

---

1. `1 / 1 point`

The Lunar Lander is a continuous state Markov Decision Process (MDP) because:

○ The reward contains numbers that are continuous valued

⦿ The state contains numbers such as position and velocity that are continuous valued.

○ The state-action value $Q(s, a)$ function outputs continuous valued numbers

○ The state has multiple numbers rather than only a single number (such as position in the x-direction)

✓ **Correct**
That's right!

---

2. `1 / 1 point`

In the learning algorithm described in the videos, we repeatedly create an artificial training set to which we apply supervised learning where the input $x = (s, a)$ and the target, constructed using Bellman's equations, is y = _____?

○ $y = R(s')$ where s' is the state you get to after taking action a in state s

○ $y = \max\limits_{a'} Q(s', a')$ where s' is the state you get to after taking action a in state s

⦿ $y = R(s) + \gamma \max\limits_{a'} Q(s', a')$ where s' is the state you get to after taking action a in state s

○ $y = R(s)$

✓ **Correct**

---

3. `1 / 1 point`

You have reached the final practice quiz of this class! What does that mean? (Please check all the answers, because all of them are correct!)

☑ The DeepLearning.AI and Stanford Online teams would like to give you a round of applause!

✓ **Correct**

☑ What an accomplishment -- you made it!

✓ **Correct**

☑ Andrew sends his heartfelt congratulations to you!

✓ **Correct**

☑ You deserve to celebrate!

✓ **Correct**